

DMA

18-1 ECE 322 Computer Organization, Lab 08

Hansung Kim
2014-16824

1 Introduction

Implement a DMA controller on top of the previous CPU implementation. Learn how CPU cooperates with DMA to achieve better performance.

2 Design

The key point of this design is the use of a signal called `d_cache_busy` that indicates whether the cache is occupying the data and address bus, i.e. handling cache miss. The addition of this single signal simplifies the logic of granting of the bus to the DMA.

2.1 Bus grant and reclaim

The CPU can grant the bus access to the DMA controller whenever the cache is not handling miss, i.e. `d_cache_busy == 0`. Conversely, if the DMA controller stops requesting for the bus, it is guaranteed that the bus is vacant and the CPU can reclaim the access immediately. Therefore, the code for the bus grant and reclaim logic is simply as follows:

```
if (bus_request && !d_cache_busy) begin
    bus_granted <= 1;
end

if (!bus_request) begin
    bus_granted <= 0;
end
```

Figure 1: Bus grant and reclaim logic using `d_cache_busy`

where `d_cache_busy` is produced by testing if the cache is being read or write but its value is not ready, i.e. `(d_readC OR d_writeC) && !d_readyC`.

2.2 Cycle stealing

The DMA cycle stealing of the CPU can be implemented easily with little modification. After each memory write of 4 words, the DMA controller sets `bus_request` to zero for exactly one cycle. This causes the CPU to retry any currently blocked memory operation in that single cycle. If there were any, the operation will set `d_cache_busy` on, delaying the `bus_granted` to go up until the operation is finished. If there were no blocked memory operation, the logic in Figure 1 would set `bus_granted` back on immediately. Therefore, the Figure 1 logic is capable of handling cycle stealing by itself.

3 Implementation

3.1 DMA module

A new module called DMA (`dma.v`) is created to implement the DMA-side transaction logic. The logic is implemented by the use of `cmd`, `BG`, `offset` and `doneM` signals.

1. Upon the CPU interrupt from the device, the CPU sets `cmd` with non-zero command signals. The DMA controller listens on this signal on each clock edge, and turns `BR` to one. It also caches the `cmd` to its internal register to keep it alive during the transaction.
2. The CPU sets `BG` to on whenever its memory operation that had been running at the point of interrupt is finished. Upon detecting this `BG`, the DMA controller issues memory write using the designated address and length in the `cmd`.
3. The controller is notified of the write finish via `doneM` signal from the memory module. Upon each `doneM`, the controller increments the `offset` counter held in an internal register. If the offset matches `len - 1`, the controller determines the data write is over, and resets both `offset` and `BR` to zero.
4. The clearing of `BR` triggers the CPU to retry its blocking memory operation and resume the pipeline. Finally, the controller sets the interrupt signal for 1 cycle and goes back to listening on `cmd`.

3.1.1 Modification for cycle stealing

To implement the cycle stealing, the step 3 of the above process should be modified to reset `BR` and set it back on immediately, for every memory write finish (`doneM`). This causes the CPU-side logic (`datapath.v`) to automatically retry its blocking memory operation, as in step 4, upon every 4-word memory write. It will subsequently set `BG` to on, causing the DMA to proceed for remaining data writes. As this process is made possible by only changing the set and reset timing of the `BR` signal, no change in the CPU-side logic is required.

3.2 Changes to the datapath

The bus granting and reclaiming logic (Figure 1), which is the CPU-side handling of the transaction, is implemented in the datapath module (`datapath.v`) with some additional logic to construct the command. Also, the stalling of the pipeline upon bus request is implemented in the hazard unit submodule (`hazard_unit.v`).

3.2.1 Command signal structure

The DMA command is simply implemented as a concatenation of the address of the memory to write and the length of the external data. In other words, it is a 2-word vector whose `cmd[31:16]` is `addr[15:0]`, and `cmd[15:0]` is `len[15:0]`.

3.2.2 Changes to the hazard unit

The hazard unit stalls the entire pipeline on `d_cache_busy`. This single check happens to be enough for handling both the D-cache miss (for ordinary load/store instructions) and the bus sharing with DMA, since `d_cache_busy` will be set to zero in both cases. Thus, the hazard unit does not require any modification from the cache implementation (aside from the variable renaming).

3.3 Changes to the cache

The cache is modified to not issue any memory read/write when the bus is granted. Its “ready” state should also remain false for the whole time span of bus grant. Thus, the `readM`, `writeM` and `readyC` signal is additionally logical-ANDed with `!bus_granted`, which effectively no-ops the cache operation when the bus is granted to the DMA.

3.4 Waveform

The waveform plot of the related signals is shown in Figure 2 and 3. Figure 2 shows the delayed bus grant due to the memory operation that was running in the CPU (`writeC` and `writeM` due to D-cache write-through) at the point of the DMA bus request. Figure 3 shows the successful cycle stealing of the CPU, where it resumes its pending memory operation in the middle of the DMA operation (`writeC` and `writeM` of the D-cache), which was triggered by the DMA temporarily disabling its bus request.

From the `pc` signal at the top, it can also be confirmed that other memory-independent operations proceed to run in the pipeline during the DMA operation.

To confirm that the DMA had actually written to the memory, the `d_dma_mem` signal at the bottom shows the memory content at the designated address (`memory[0x01f4]`), which indeed changes to the `edata` value immediately after the first DMA write.

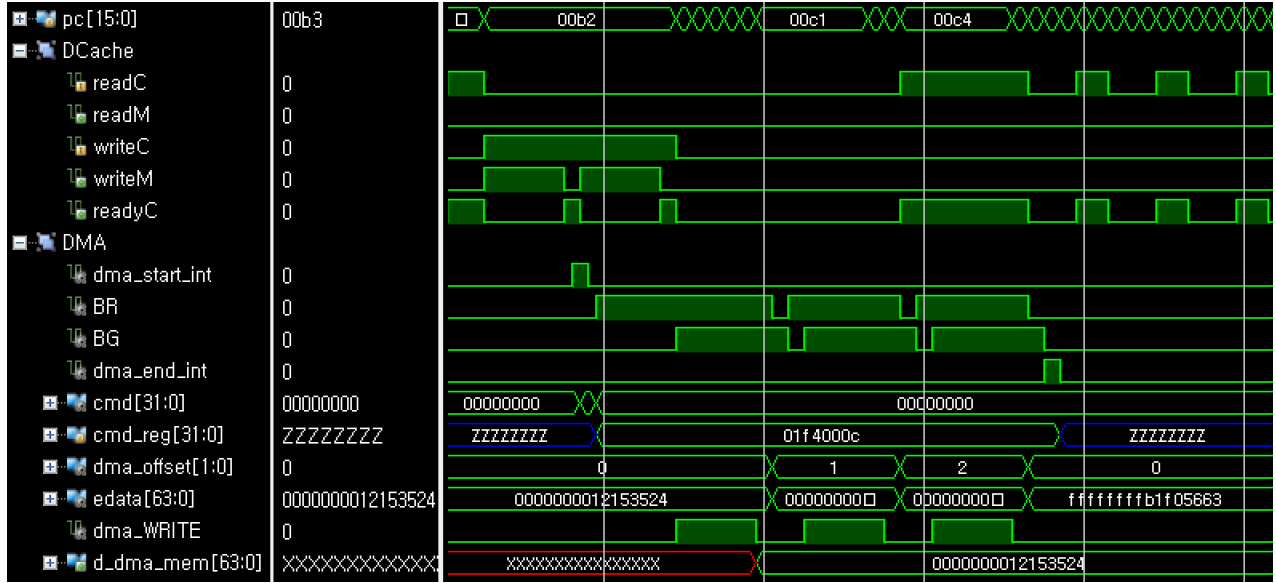


Figure 2: Waveform of the DMA and D-cache signals with no successful cycle stealing




Figure 3: Waveform of the DMA and D-cache signals with two successful cycle stealing

4 Discussion

A simple DMA controller along with cycle stealing could be successfully implemented. Compared to the previous version, the number of clock cycles to finish the testbench increased from 2796 to 2802, which is natural considering the additional latency caused by DMA occupying the memory bus and preventing the pipeline to proceed.

A quick test showed that the implementation without cycle stealing shows the clock cycles of 2806. The performance increase due to cycle stealing is marginal. However, cycle stealing has the

advantage that it does not leave the CPU completely unresponsive during the device communication, which could be a major problem for slow devices that require long I/O time.

5 Conclusion

Successfully implemented a DMA controller on top of the previous CPU implementation. Learned how CPU cooperates with DMA to achieve better performance.